

Agent skill learning and keepaway using parameterized policy search

By Rahul Dass

Advisor: Dr. Ubbo Visser

March 26, 2018

Outline

1. My interests: Reinforcement Learning (RL)
 - a. RL in a nutshell
 - b. MDP: planning and learning
 - c. RL families: MFL and MBL
2. RL case study: RoboCup 3D Simulation Domain
 - a. Layered Learning (UT Austin)
3. RL application: RoboCanes 3D Simulation Team
 - a. How can RL be applied?
 - i. Robot skill learning
 - ii. Keepaway planning

1. RL in a nutshell

- Let x : instance (featurized representation)
 - What type of robot? Number of D.O.F? What kind of skills?
- Let a : action taken
 - Sprinted towards a soccer ball
- Let v : evaluative feedback
 - +2 for getting to the ball; -5 for falling down
- **Goal: construct a decision rule that maximizes expected value^[3, 4]**

Why is this notion exciting?

Advantages

- Don't tell the learner **what to do**
- Learner just needs to know **how to score**

Challenges

- Nature of evaluation is **weak!**
- Receiving scalar values alone, agent **does not know** what it should be doing
- Agent cannot tell if the scalar feedback is a **reward (+ve)** or **penalty (-ve)**

RL feedback signal

- **Evaluative**

- Not telling a robot **which way to turn**. Just when it sees the ball - **signal to confirm or not**.

- **Sampled**

- Input that AIBOS gets is a colored histogram coming from cameras. Probably will not see the same colored patterns twice. Has to **generalize** between color patterns.

- **Sequential**

- Robot starting in some initial s , takes a , gets **no immediate v** due to a taken. It continues making (s, a) pair movements.
- When ball is located - it was somehow **set up** by the **sequence of a 's taken**

MDP: planning and learning

Q: How do we make agents take decisions in order to maximize reward?

- **Markov Decision Process** contains elements; the environment has 2-functions:
 - States, actions/decisions (discrete)
 - Transitions, rewards stationary, Markovian
- Transition function: $Pr(s' | s, a) = T(s, a, s')$
 - T takes current (s, a) pair
 - Output: **probability distribution** over all next states s'
- Reward function: $E(r | s', a) = R(s, a)$
 - Mapping agent being in s and taking a
 - E provides an evaluation of **how good** that a was

Agent's perspective

Optimal policy: **which actions** to take for **which states** that **maximizes** the total cumulative (expected, discounted) reward

- **Policy**: $\pi^*(s) = \operatorname{argmax}_a Q^*(s, a)$
 - Mapping of states \rightarrow actions
- $Q^*(s, a) = R(s, a) + \gamma \sum_i T(s, a, s') \max_{a'} Q^*(s', a')$
 - Assume from s' till end of trajectory, agent behaves optimally
 - Q^* is an estimate of the **total expected, discounted future reward** for a single-step transition
 - This is the Bellman equation

Solving the Bellman equation = Decision making in a MDP

Solving Bellman: upshot

Given a reasonably sized MDP, there are 3 algorithms that can figure out the optimal way for an agent to behave that maximizes expected reward:

1. Value iteration converges in limit
2. Policy iteration converges in finite time
3. Linear programming runs in polynomial time

Downside:

- Major assumption that we'll be given such a MDP
- **RL problem**: agent is living in a world, and it needs to figure out **on its own** how the world works etc.^[3]

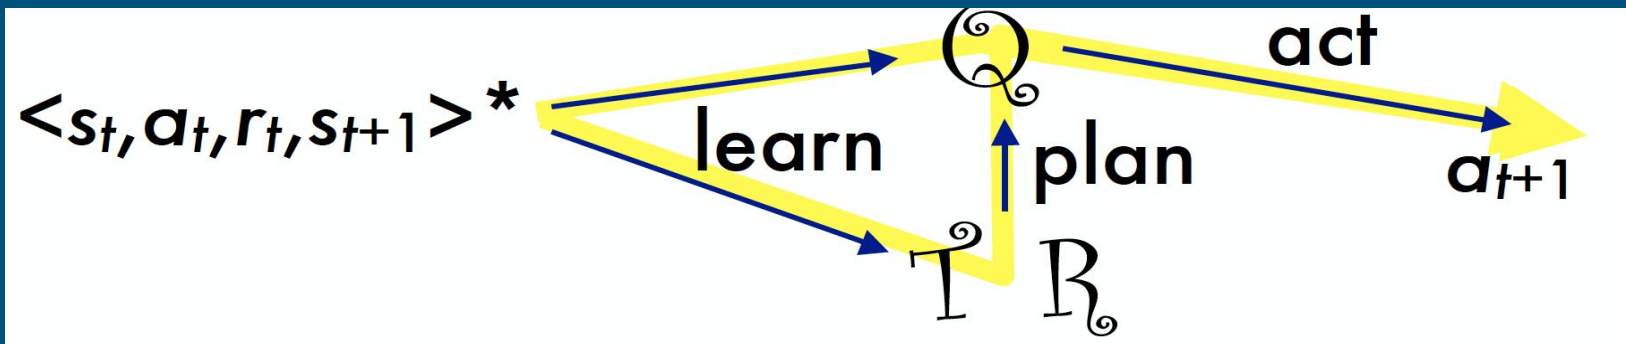
RL families^[3]

Model-free learning

- Use experience $\langle s, a, r, s' \rangle$ to “learn” Q directly, then be able to take a
- Eg: Q-learning and SARSA

Model-based learning

- Use experience $\langle s, a, r, s' \rangle$ to “learn” T and R , which are used to compute Q and then can the agent determine what a to take
- Eg: R_{MAX} and E^3



2. RL case study

RoboCup 3D Simulation Domain^[5, 6]

- Teams of 11 vs 11 autonomous robots playing soccer
- Realistic physics using Open Dynamics Engine (ODE)
- Simulated robots are modeled after the Aldebaran Nao robot
- Robot receives noisy visual information about environment
- Robots can communicate with each other over limited bandwidth channel

Head-to-Head

UT Austin Villa^[6]

- **2007-2009:** no goals scored, no games won
- **2010:** improved walk but still unstable and falling frequently
- **2011-2017:** current champion
 - RL based **Layered Learning** approach

RoboCanes

- **Champion:** German Open (2011), US Open (2015)
- **2nd Place:** World Cup (2012, 2014), Iran Open (2011, 2012), US Open (2014)
- **3rd Place:** Iran Open (2014)
- **4th Place:** World Cup (2017)
- **Quarterfinalist:** World Cup (2015, 2013, 2011, 2010)

Direct policy search^[3, 4]

- **Goal:** learn a **parameterized policy** that determines an agent's behavior
- **Method:** Optimization algorithm that produces **candidate parameters** for an agent to evaluate an optimization task, eg: robot kicking a ball
- **Evaluation and outcome:** upon task completion, agent **evaluates parameters** and **returns fitness** (reward)
- **Repeat:** algorithm generates **new parameters** to find a (optimal) policy that hopefully **improves** robot's ability and performance leading to **higher fitness**

CMA-ES as a baseline

Covariance **M**atrix **A**daptation **E**volutionary **S**trategy^[13] is an evolutionary numerical optimization method

- Similar to a genetic algorithm being **generation** and **population based**
- However, unlike a genetic algorithm that uses a genetic operator, CMA-ES **samples** candidate members / **parameter sets** from a **multivariate Gaussian distribution**
- One generation of learning to the next, CMA-ES adjust the distribution to move to areas of higher fitness in parameter search space^[5, 6]

Layered Learning^[7]

- Hierarchical ML paradigm enabling learning of complex behaviors by **incrementally learning a series of sub-behaviors**
 - Higher layers directly depend on the **learned lower layers**

Goal: create “overall” optimal behavior policy

1. Sequential Layered Learning^[7]

- a. After **parameters** are learned for **first layer**, they are **frozen**, then learn parameters for next layer
- b. Problem: **Too limiting** in a **joint behavior policy search space**. Parameters from first layer is cut off from next layer.

2. Concurrent Layered Learning^[8]

- a. Learning parameters from first layer are **kept open**, when parameters for next layer are learning
- b. Problem: **Increased dimensionality** causes learning to be **too hard / slow** or **intractable**

Proposal: Overlapping Layered Learning^[9]

Layered learning, both in **series** and **parallel**, where **some parameters** of previously learned layers are left open during learning of subsequent layers

- Tradeoff: freezing and keeping parameters open (overcomes SLL and CLL)
- Optimizes overlap between behaviors

1. Combining Independently Learned Behaviors (CILB)

- a. Learning 2 or more behaviors in **parallel**, relearn **subset** in next layer (conflicting behaviors)

2. Partial Concurrent Layered Learning (PCLL)

- a. Parameters are **partly open** after first layer of learning (highlighting tradoff)

3. Previously Learned Layer Refinement (PLLR)

- a. After behavior is learned + params frozen for two layers, part of first layer opened to relearn in presence of latter layer

3. RL application: RoboCanes 3D Sim Team

- What are we able to do so far?
 - Learning parameters from the walk engine
 - Learning of the kick is via CMA-ES, running
- Agent skill learning:
 - **Task 1:** Find an optimal approach to the ball via parameterized policy search
 - Challenge: Without a stable walk engine, policy search parameters will get affected
 - **Task 2:** Design a new walk based off double inverted pendulum approach (UT Austin)
 - Preliminaries: Cartpole toy problem OpenAI gym
 - **Task 3:** Multi-skill learning
 - Stable transitioning between skills, say: walk \rightarrow kick \rightarrow walk
 - Challenges: delay in ball positioning and robot instability

Keepaway planning

- Start with simplest 3 vs 2 keepaway scenario^[10]
 - 3 agents from same team keeping possession of the ball from 2 agents attempting to take the ball away
 - Preliminaries: simplify problem in a 2D setting, modeled using the classic RL gridworld problem: agent getting to target with obstacles
 - Challenges: multi-agent skill learning and movement coordination

Thank you!

Dr. Visser - giving me the opportunity to do my Ph.D. here at UM. Being my mentor. Introducing me to the world of RL and robotics.

Dr. Schwartz - constant source of advice and support, especially our discussions of *RL in the brain, and Deep Learning*.

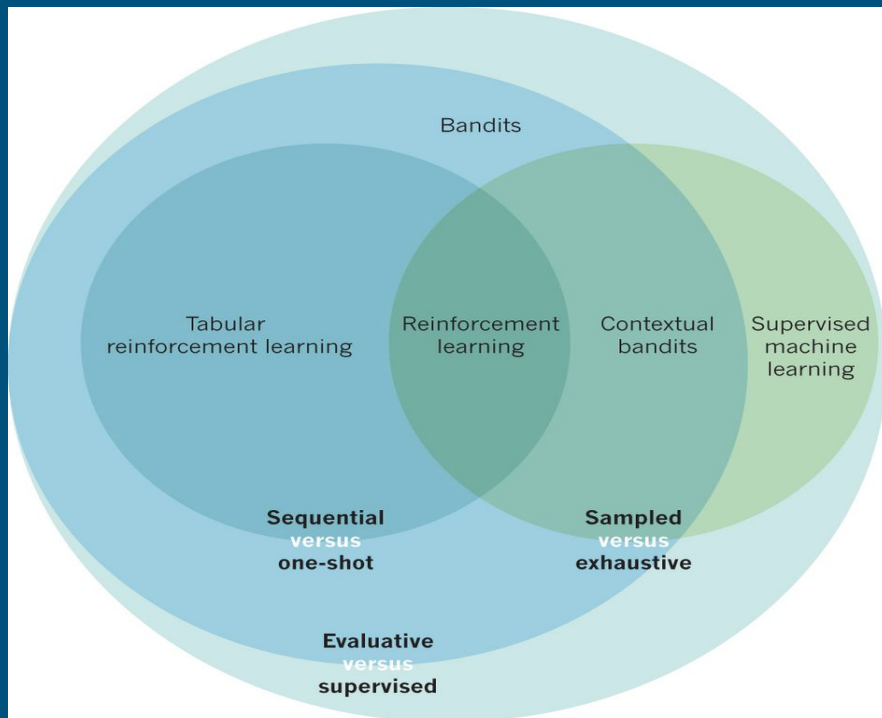
Dr. Sutcliffe - teaching me *how to give a talk!*

RoboCanes family - past, present and future

References

- [1] Littman M., *Reinforcement Learning improves behavior from evaluative feedback* (**Nature**, 2015)
- [2] Sutton, R. and Barto A. G., *Reinforcement Learning: An Introduction 1st Edition* (**MIT Press**, 1998)
- [3] Littman M., *Basics of Computational Reinforcement Learning* (**RLDM**, 2015)
- [4] Silver D., *Advanced Topics: Reinforcement Learning Course* (**UCL**, 2015)
- [5] Stone P., *Robotic Skill Learning: From Real World to Simulation and Back - CMU RI seminar* (2017)
- [6] MacAlpine P., *Multilayered Skill Learning and Movement Coordination for Autonomous Robotics Agents - Ph.D. Defense* (2017)
- [7] Stone, P., and Veloso M., *Layered Learning* (**Springer**, 2000)
- [8] Whiteson S., and Stone P., *Concurrent Layered Learning* (**AAMAS**, 2003)
- [9] MacAlpine P., and Stone P., *Overlapping Layered Learning* (**AIJ**, 2018)
- [10] Stone P., Sutton R., and Singh S., *Reinforcement Learning for 3 vs 2 keepaway* (2001)
- [11] Hansen, N., Muller S. D., and Koumoutsakos P., *Reducing the Time Complexity of Derandomized Evolutionary Strategy with Covariance Matrix Adaptation (CMA-ES)* (2003)

Feedback revisited^[1]



...there's quite a lot!

Solving Bellman: 3 algorithms^[1-4]

a. Value iteration converges in limit

- i. Guess $Q_0(s, a)$. Substitute in Bellman to compute new Q -values by taking max (nonlinear). Q -value. Keep iterating, in the limit, converges to a Q -value that is solution to Bellman.

b. Policy iteration converges in finite time

- i. Guess $Q_0(s, a)$. Use it to build a policy $\pi_t(s)$ where a taken gives highest estimate of Q -value. Evaluate π_t using Bellman using Q . Either it returns same π as before (optimal), or gives a better one from which we generate new Q -value etc. Linear constraints \rightarrow finite time.

c. Linear programming runs in polynomial time

- i. Express (if possible) MDP as a linear program. Solve it. Solution to Bellman found